

ML_project_template__

March 29, 2023

1 data Gathering & preprocessing

Loading data from a file:

```
[ ]: import pandas as pd

data = pd.read_csv('data.csv')
```

Handling missing data:

```
[ ]: # drop rows with missing values
data.dropna(inplace=True)

# fill missing values with the mean of the column
data.fillna(data.mean(), inplace=True)
```

Encoding categorical variables:

```
[ ]: # create dummy variables for categorical variables
data = pd.get_dummies(data, columns=['color'])

# label encoding for ordinal categorical variables
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
data['size'] = le.fit_transform(data['size'])
```

Scaling numerical variables:

```
[ ]: # standardize data to have mean=0 and variance=1
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
data[['age', 'income']] = scaler.fit_transform(data[['age', 'income']])

# normalize data to have values between 0 and 1
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
data[['age', 'income']] = scaler.fit_transform(data[['age', 'income']])
```

Handling outliers:

```
[ ]: # remove outliers using z-score
from scipy import stats
data = data[(np.abs(stats.zscore(data)) < 3).all(axis=1)]

# winsorize outliers to a specified percentile
from scipy.stats.mstats import winsorize
data['income'] = winsorize(data['income'], limits=[0.05, 0.05])
```

2 Data wrangling

Data wrangling, also known as data munging, is the process of cleaning, transforming, and reshaping data into a format that is suitable for analysis. Here are some common data wrangling steps and corresponding Python codes using the Pandas library:

renaming columns

```
[ ]: # rename a single column
data.rename(columns={'old_name': 'new_name'}, inplace=True)

# rename multiple columns
data.rename(columns={'old_name1': 'new_name1', 'old_name2': 'new_name2'},
            inplace=True)
```

Removing duplicates:

```
[ ]: import pandas as pd

data = pd.read_csv('data.csv')

# remove duplicates based on all columns
data.drop_duplicates(inplace=True)

# remove duplicates based on a specific column
data.drop_duplicates(subset=['column_name'], inplace=True)
```

Filtering rows:

```
[ ]: # filter rows based on a condition
filtered_data = data[data['age'] > 25]

# filter rows based on multiple conditions
filtered_data = data[(data['age'] > 25) & (data['income'] > 50000)]

# filter rows based on a list of values
filtered_data = data[data['color'].isin(['red', 'green', 'blue'])]
```

Handling missing data:

```
[ ]: # drop rows with missing values
data.dropna(inplace=True)

# fill missing values with the mean of the column
data.fillna(data.mean(), inplace=True)
```

Aggregating data:

```
[ ]: # group data by a categorical variable and compute mean for each group
grouped_data = data.groupby(['category'])['value'].mean().reset_index()

# group data by a categorical variable and compute multiple summary statistics
↳ for each group
grouped_data = data.groupby(['category']).agg({'value': ['mean', 'median', '
↳ max', 'min']}).reset_index()
```

Pivot tables:

```
[ ]: # create a pivot table to summarize data by multiple variables
pivot_table = data.pivot_table(index=['category'], columns=['year'],
↳ values='value', aggfunc='sum')
```

3 Data analysis

```
[ ]: Data analysis refers to the process of using statistical methods and models to
↳ gain insights and draw conclusions from data.
Here are some common data analysis steps and their corresponding Python code
using various libraries such as Pandas, NumPy, Matplotlib, and Scikit-learn:
```

Descriptive statistics:

```
[ ]: import pandas as pd

data = pd.read_csv('data.csv')

# calculate summary statistics
summary_stats = data.describe()

# calculate correlation matrix
corr_matrix = data.corr()
```

Data visualization:

```
[ ]: import matplotlib.pyplot as plt

# create a scatter plot
plt.scatter(data['age'], data['income'])
plt.xlabel('Age')
```

```

plt.ylabel('Income')
plt.title('Age vs Income')

# create a histogram
plt.hist(data['age'], bins=20)
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.title('Age Distribution')

# create a bar chart
counts = data['category'].value_counts()
plt.bar(counts.index, counts.values)
plt.xlabel('Category')
plt.ylabel('Count')
plt.title('Category Counts')

```

Hypothesis testing:

```

[ ]: from scipy.stats import ttest_ind

# perform a t-test for two independent samples
group1 = data[data['category'] == 'A']['value']
group2 = data[data['category'] == 'B']['value']
t_stat, p_value = ttest_ind(group1, group2)

```

Machine learning modeling:

```

[ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

data = pd.read_csv('data.csv')

# split data into training and test sets
X = data[['age', 'income']]
y = data['value']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    random_state=42)

# train a linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# make predictions on the test set
y_pred = model.predict(X_test)

# evaluate model performance using mean squared error

```

```
mse = mean_squared_error(y_test, y_pred)
```

4 Training a model

Training a model involves using a machine learning algorithm to learn patterns and relationships in the data, and create a model that can make predictions on new data. Here are the general steps for training a model and their corresponding Python code using Scikit-learn:

Splitting data into training and test sets:

```
[ ]: import pandas as pd
      from sklearn.model_selection import train_test_split

      data = pd.read_csv('data.csv')

      # split data into training and test sets
      X = data[['feature1', 'feature2', ...]]
      y = data['target']
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      ↪random_state=42)
```

Preprocessing the data:

```
[ ]: from sklearn.preprocessing import StandardScaler

      # standardize the training data
      scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)

      # apply the same transformation to the test data
      X_test = scaler.transform(X_test)
```

Choosing and training a model:

```
[ ]: from sklearn.linear_model import LinearRegression

      # initialize a linear regression model
      model = LinearRegression()

      # train the model on the training data
      model.fit(X_train, y_train)
```

Evaluating the model on the test set:

```
[ ]: from sklearn.metrics import mean_squared_error

      # make predictions on the test set
      y_pred = model.predict(X_test)
```

```
# evaluate model performance using mean squared error
mse = mean_squared_error(y_test, y_pred)
```

Tuning the model hyperparameters:

```
[ ]: from sklearn.model_selection import GridSearchCV

# define a grid of hyperparameters to search over
param_grid = {'alpha': [0.01, 0.1, 1.0, 10.0]}

# perform a grid search to find the best hyperparameters
grid_search = GridSearchCV(model, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# use the best hyperparameters to train the final model
model = grid_search.best_estimator_
model.fit(X_train, y_train)
```

5 Testing a trained model

Testing a trained model involves using it to make predictions on a new dataset and evaluating its performance. Here are the general steps for testing a model and their corresponding Python code using Scikit-learn:

Loading the test data:

```
[ ]: import pandas as pd

test_data = pd.read_csv('test_data.csv')
```

Preprocessing the test data:

```
[ ]: from sklearn.preprocessing import StandardScaler

# standardize the test data using the same scaler as the training data
X_test = scaler.transform(test_data[['feature1', 'feature2', ...]])
```

Making predictions with the trained model:

```
[ ]: # make predictions on the test data using the trained model
y_pred = model.predict(X_test)
```

Evaluating the model's performance on the test set:

```
[ ]: from sklearn.metrics import mean_squared_error

# evaluate model performance using mean squared error
mse = mean_squared_error(test_data['target'], y_pred)
```

Visualizing the model's performance:

```
[ ]: import matplotlib.pyplot as plt

# create a scatter plot of predicted vs actual values
plt.scatter(test_data['target'], y_pred)
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Predicted vs Actual Values')
```